

Today's section will quickly debrief Problem Set 2 and introduce spatial data analysis in R.

## Problem Set Debrief

- Many people didn't realize that the correlation coefficient between  $\ln(Q)$  and the residuals is zero by construction in OLS.
- $J = 8$  in the F-test as each extra group coefficient must be set to zero individually in the restricted model. 4 more groups than in Q2 so 4 slopes + 4 intercepts = 8.
- Some people calculated the F-test the hard way. See the solutions.
- You can restrict the number of digits and add commas to numbers in knitr using the `prettyNum()` function.
- Lack of comments on your submissions indicates that things look mostly correct.

## Spatial Data Analysis

R is the first (and mostly only) language with which I have performed spatial data analysis. It's not a nice, coherent, well formed, piece of software for this purpose. Of the packages in my package install script (which you should all have for when you need to install/reinstall R), I counted 18 items with functions related to spatial analysis or mapping. It's a mess. Before the `sp` package, apparently it used to be worse! When I'm a tenured professor who has run out of ideas, or an unemployed bum living at my mother's house, I plan to start a project that collects the best of all these functions into a single package with a text book. In the meantime, the status quo will have to do. A decent resource is Applied Spatial Data Analysis with R (second edition) by Bivand, Pebesma, and Gómez-Rubio.

Why should I use R for spatial analysis then? Two reasons: 1) the value of keeping your work in one program for ease of reproducibility and 2) active development and voluntary support. People are keenly working on developing R's spatial capabilities and people answer questions. The R-sig-geo mailing list is a good resource, as is stackoverflow.

## Spatial Data Types

The first step in understanding spatial analysis in R is understanding the various data types. This is going to be quick. For each of these classes, there is a corresponding class that adds a `data.frame` with a row for each element of the class, so that data on each element can be nicely kept with the spatial object. For `SpatialPoints`, for example, the corresponding object that also contains a `data.frame` is the `SpatialPointsDataFrame`.

### SpatialPoints

A `SpatialPoints` object is a collection of points represented by coordinates, usually latitude/longitude pairs. Each element is a single point.

## SpatialLines

A `SpatialLines` object is a collection of lines represented by collections of coordinates in an order. Each element is a collection of points, ordered to represent the line.

## SpatialPolygons

A `SpatialPolygons` object is simply a collection of closed lines. Individual “polygons” can have holes, and comprise more than one actual `Polygon` object. Mostly, you won’t have to concern yourself with the fine details of how these objects work.

## SpatialGrid/SpatialPixels

`SpatialGrid/SpatialPixels` are basically the same object and represent a something close to a rectangular grid. These are used to represent, for example, satellite imagery and gridded weather.

## Reading spatial data

There are many different file types that represent spatial data and I have not yet come across a file type that can’t be read, though finding a method that works is not always easy. The GDAL driver is pretty flexible and can handle a lot of formats. We’ll look at a couple:

## ESRI Shapefiles

Being a popular format, this one is particularly easy to get into R. Let’s download a USA states shapefile and read it in. You will need to install the `sp` and `rgdal` packages if you haven’t already.

```
library(sp)
library(rgdal)

## Loading required package: methods
## rgdal: version: 0.9-1, (SVN revision 518)
## Geospatial Data Abstraction Library extensions to R successfully loaded
## Loaded GDAL runtime: GDAL 1.11.0, released 2014/04/16
## Path to GDAL shared files: C:/Users/Kenny/R/win-library/3.1/rgdal/gdal
## GDAL does not use iconv for recoding strings.
## Loaded PROJ.4 runtime: Rel. 4.8.0, 6 March 2012, [PJ_VERSION: 480]
## Path to PROJ.4 shared files: C:/Users/Kenny/R/win-library/3.1/rgdal/proj

# Download a USA shapefile
tmpFile <- tempfile()
download.file("http://www2.census.gov/geo/tiger/GENZ2013/cb_2013_us_state_20m.zip",
             destfile = tmpFile)
unzipped <- unzip(zipfile = tmpFile, exdir = getwd())

# This is the actual reading stage
```

```

states <- readOGR(dsn = getwd(), layer = "cb_2013_us_state_20m")

## OGR data source with driver: ESRI Shapefile
## Source: "C:/Users/Kenny/Dropbox/PhD/ARE212GSI/Github/ARE212/section-11", layer: "cb_2013_us_s
## with 52 features and 9 fields
## Feature type: wkbPolygon with 2 dimensions

# Remove all the files on disk
file.remove(tmpFile)

## [1] TRUE

file.remove(unzipped)

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE

```

Now let's take a quick look at what we've downloaded!

```
plot(states)
```



Could look better - but we'll worry about that more next time. We can certainly confirm that we've downloaded a `SpatialPolygons` object for the USA.

To look closer at what we have, let's check out the slots<sup>1</sup> that belong to the `SpatialPolygonsDataFrame` class.

```
slotNames(states)

## [1] "data"          "polygons"      "plotOrder"     "bbox"          "proj4string"
```

The `data` slot simply contains a `data.frame` with a row for each `Polygons` object, here one for each

<sup>1</sup>Think of slots as attributes for a given class that are defined for all objects that belong to that class.

state. The `polygons` slot contains a list of `Polygons` objects. You may interact with the rest at some stage, but I wouldn't worry about them for now. Let's look at the `data.frame` part.

```
head(states@data)
```

##	STATEFP	STATENS	AFFGEOID	GEOID	STUSPS	NAME	LSAD	ALAND
## 0	01	01779775	0400000US01	01	AL	Alabama	00 131172434095	
## 1	05	00068085	0400000US05	05	AR	Arkansas	00 134772954601	
## 2	06	01779778	0400000US06	06	CA	California	00 403482685922	
## 3	09	01779780	0400000US09	09	CT	Connecticut	00 12541965607	
## 4	12	00294478	0400000US12	12	FL	Florida	00 138897453172	
## 5	13	01705317	0400000US13	13	GA	Georgia	00 148962779995	
##	AWATER							
## 0	4594920201							
## 1	2958815561							
## 2	20484304865							
## 3	1815409624							
## 4	31413676956							
## 5	4947803555							

### .bil pixel files

Now, let's download some temperature data as a `SpatialGridDataFrame`. I'll download the maximum temperature for the contiguous USA for my birthday last year. You can see that now we use the `readGDAL()` function, rather than the `readOGR()` function. Google is your friend for working out when to use which reading function to use to get the data in. Always favor solutions that use the `rgdal` package, as these functions are able to read the map projections, while others do not (for example, the `maptools` package).

```
tmpFile <- tempfile()

url <- paste0("http://www.prism.oregonstate.edu/fetchData.php?type",
              "=bil&kind=recent&elem=tmax&range=daily&temporal=",
              "20140125")

# Download the file to fileName
download.file(url = url, mode = "wb", destfile = tmpFile)


# unzip the file to the working directory.
unzippedFiles <- unzip(zipfile = tmpFile, overwrite = TRUE)

# read into R
tMaxGrid <- readGDAL(fname = unzippedFiles[1])

## ./PRISM_tmax_stable_4kmD1_20140125_bil.bil has GDAL driver EHdr
## and has 621 rows and 1405 columns

# delete the file
file.remove(tmpFile)
```

```
## [1] TRUE
file.remove(unzippedFiles)
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
image(tMaxGrid)
```



Great! This seems to have worked.

## Google maps distances

Now, let's use R's interface with Google maps to get some distances.

```
library(ggmap)
distances <- mapdist(from = "Giannini Hall, UC Berkeley, CA 94720",
                     to = as.character(states@data$NAME),
                     mode = "driving") # $
states$distanceToGiannini <- distances$km
```

## Spatial Data Analysis, finally

Now, let's quickly do something silly and see if there is a significant correlation between driving distance to Giannini and the temperature grid you downloaded.

```
# This gets the simple average of the grid cells whose centres
# are in each state. NOT a good solution for smaller polygons
# like ZIPs. Stored in a SpatialPolygonsDataFrame.
tMax <- aggregate(tMaxGrid, by = states, FUN = mean, na.rm=TRUE)

# Access the temperature and store it in the states object.
states$tMax <- tMax@data$band1
```

Now, we're ready to run a regression!

```
summary(lm(distanceToGiannini ~ tMax, data = states@data))

##
## Call:
## lm(formula = distanceToGiannini ~ tMax, data = states@data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1645.48  -446.80   -31.31   358.87  2967.96
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3552.67     112.97   31.448  < 2e-16 ***
## tMax        -148.09      15.91   -9.307 3.09e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 780.2 on 47 degrees of freedom
## (3 observations deleted due to missingness)
## Multiple R-squared:  0.6482, Adjusted R-squared:  0.6408
## F-statistic: 86.61 on 1 and 47 DF,  p-value: 3.09e-12
```

Obviously, there is a lot more to spatial analysis in R than calculating attribute values and running regressions. We can calculate spatial correlograms, calculate network statistics, all sorts. Mostly you will discover these things along the way - I'm always available for conversations about these things to get you on the right track.